

# Lecture 8a: 2-3-4 and Red-Black Trees

---

CS 70: Data Structures and Program Development

Tuesday, March 10, 2020

# The Story So Far

Lookup and insert with  $n$  nodes:

- naive BST: worst-case  $\Theta(n)$
- randomized BST: expected  $\Theta(\log n)$  and worst-case  $\Theta(n)$

Future preview:

- splay tree: *amortized*  $\Theta(\log n)$  and worst-case  $\Theta(n)$ .

How can we get worst-case  $\Theta(\log n)$ ?

- Need to ensure the tree has height  $\Theta(\log n)$ ?


# A Perfect Tree

All levels are full.

Has  $2^{h+1} - 1$  nodes, where  $h$  is height.

So  $h \in \Theta(\log n)$ .

This is a lower bound! For any BST,  $h \in \Omega(\log n)$ .



# Red-Black Trees



# Red-Black Trees (definition 1)

A red-black tree is a BST such that:

- Every node is red or black
- The root is black
- No red parent has a red child
- Every path from the root to `nullptr` passes through the same number of black nodes.

Claim: an  $n$ -node red-black tree has height  $\Theta(\log n)$

# Valid Red-Black Trees?



## Red-Black Trees have $\Theta(\log n)$ height.

Suppose every path to `nullptr` hits  $b$  black nodes.

1. The height is less than  $2b$ .
2. The tree has at least  $2^b - 1$  nodes, i.e.,  $2^b - 1 \leq n$ .

# Red-Black Trees have $\Theta(\log n)$ height.

Suppose every path to `nullptr` hits  $b$  black nodes.

1. The height is less than  $2b$ .
2. The tree has at least  $2^b - 1$  nodes, i.e.,  $2^b - 1 \leq n$ .
3.  $h < 2b \leq 2 \log_2(n + 1)$ , so  $h \in O(\log n)$ .



# Red-Black Trees have $\Theta(\log n)$ height.

Suppose every path to `nullptr` hits  $b$  black nodes.

1. The height is less than  $2b$ .
2. The tree has at least  $2^b - 1$  nodes, i.e.,  $2^b - 1 \leq n$ .
3.  $h < 2b \leq 2 \log_2(n + 1)$ , so  $h \in O(\log n)$ .

Also we know that  $h \in \Omega(\log n)$ .

# Red-Black Trees have $\Theta(\log n)$ height.

Suppose every path to `nullptr` hits  $b$  black nodes.

1. The height is less than  $2b$ .
2. The tree has at least  $2^b - 1$  nodes, i.e.,  $2^b - 1 \leq n$ .
3.  $h < 2b \leq 2 \log_2(n + 1)$ , so  $h \in O(\log n)$ .

Also we know that  $h \in \Omega(\log n)$ .

Since  $n \in O(\log n)$  and  $n \in \Omega(\log n)$ ,  $n \in \Theta(\log n)$ .

# The Problem

Red-Black Trees can be a pain to implement.

- Lookup is easy - same as in a BST
- But how to update the colors when we insert?
  - Not immediately obvious. . .


Maybe we can find a balanced tree that is more intuitive?

# 2-3-4 Trees

# Why not require perfection at all times?


All levels are full.

Has  $2^{h+1} - 1$  nodes, where  $h$  is height.




## 2-3-4 Tree


1. Allow a node to store more than one key.
2. All leaves at the same depth



*2-node*




*3-node*



*4-node*

# Typical 2-3-4 tree



## 2-3-4 trees: insertion

### Rules for insert-as-leaf

- Always keep the leaves at the same level
- If leaf is a 2 node: Make it a 3 node
- If leaf is a 3 node: Make it a 4 node
- If we encounter a 4 node on the way down, “promote” the middle (why)?

Exercise: insert 1,2,3,4,5,6,...



## Your Turn: 2-3-4 Trees

- In groups of 4
- Starting from an empty tree
- Insert the month numbers and day numbers of your birthdays (without repeats) into a 2-3-4 Tree.

# Advantages and Disadvantages


## Advantages of 2-3-4 Trees

- The tree is always “balanced”
- Worst case for insert and lookup is  $\Theta(\log n)$  for a tree with  $n$  values
- Simple algorithm; no rotations required.
- Smaller height than a typical binary tree. (why?)


## Disadvantages?

# Encoding 2-3-4 Trees in Binary



# Super Links!




## 3-Nodes with Super links




## 4-nodes with Super links



## Exercise: Convert to Binary + Superlinks




# Coloring *pointers* is weird...






# Coloring *pointers* is weird...




Idea: color the node, not the (incoming) edge!


# Coloring *pointers* is weird...



Idea: color the node, not the (incoming) edge!



## Exercise: Convert to Binary Tree



# Word of Warning

- Red-Black Trees are efficient and commonly used (e.g., `std::set`)
- But it's easy to write convoluted, messy, confusing, scary Red-Black Tree implementations!
- If you ever implement a red-black tree, refer back to the 2-3-4 implementation.

# Advantages & Disadvantages


## Advantages of Red-Black Trees

- All the balance properties of 2-3-4 Trees
- Only one extra bit of color per node needed over a standard BST
- Only one node type, so easier to implement than 2-3-4 Trees

## Disadvantages of Red-Black Trees


- Can be messy if you don't draw diagrams and think about the 2-3-4 equivalents during implementation.
- Rotations trickier conceptually than 2-3-4 tree operations.

## Promote the 3



1. Draw the Red-Black equivalent before promotion
2. Do the promotion on the original 2-3-4 tree
3. Draw the equivalent Red-Black tree after the promotion

## Different scenario: Promote the 3



1. Draw the Red-Black equivalent before promotion
2. Do the promotion on the original 2-3-4 tree
3. Draw the equivalent Red-Black tree after the promotion

# Learning Targets

1. I can explain the fundamental idea behind Red-Black trees.
2. I can explain the fundamental idea behind 2-3-4 trees.
3. I can explain the fundamental idea behind Red-Black trees.